


Mars Pathfinder Incident

- Landing on July 4, 1997
- "...experiences software glitches..."
- Pathfinder experiences repeated RESETs after starting gathering of meteorological data.

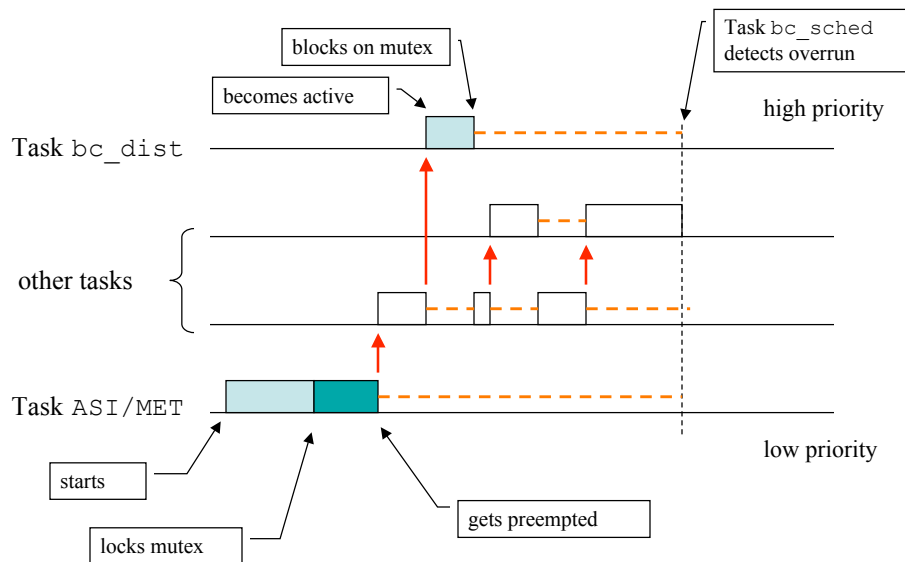
- RESETs generated by watchdog process.
- Timing overruns caused by priority inversion.

- Resources:

research.microsoft.com/~mbj/Mars_Pathfinder/Mars_Pathfinder.html



Priority Inversion on Mars Pathfinder



Mars Pathfinder: Resolution


- “Faster, better, cheaper” had NASA and JPL using “shrink-wrap” hardware (IBM RS6000) and software (Wind River VxWorks RTOS).
- Logging designed into VxWorks enabled NASA and Wind River to reproduce the failure on Earth. This reproduction made the priority inversion obvious.
- NASA patched the lander’s software to enable priority inheritance.

Resource Access

- Processor(s)
 - m types of serially reusable resources R_1, \dots, R_m
 - An execution of a job J_i requires:
 - a processor for e_i units of time
 - some resources for exclusive use
- Resources
 - serially Reusable: Allocated to one job at a time. Once allocated, held by the job until no longer needed.
 - examples: semaphores, locks, servers, ...
 - operations:
`lock(Ri) -----<critical section>----- unlock(Ri)`
 - resources allocated non-preemptively
 - critical sections properly nested

Preemption During Critical Sections


Example:



- Negative effect on schedulability and predictability.


Unpredictability: Scheduling Anomalies

- Example: $T_1 = (c_1=2, e_1=5, p_1=8)$ $T_2 = (4, 7, 22)$ $T_3 = (4, 6, 26)$




- Shorten critical section of T_3 :

$$T_1 = (c_1=2, e_1=5, p_1=8) \quad T_2 = (4, 7, 22) \quad T_3 = (2.5, 6, 26)$$



Disallow Process Preemption in CS



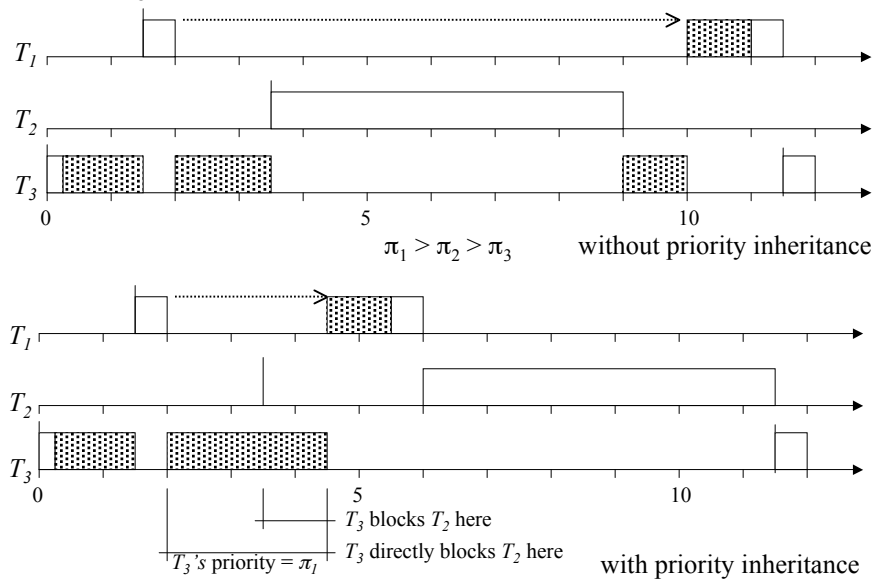
- Analysis identical to analysis with non-preemptable portions
- Define: β = maximum duration of all critical sections
- Task T_i is schedulable if

$$\sum_{k=1} \frac{e_k}{p_k} + \frac{\beta}{p_i} = U_X(i)$$

X : scheduling algorithm

- Problem: critical sections can be rather long.

Priority Inheritance



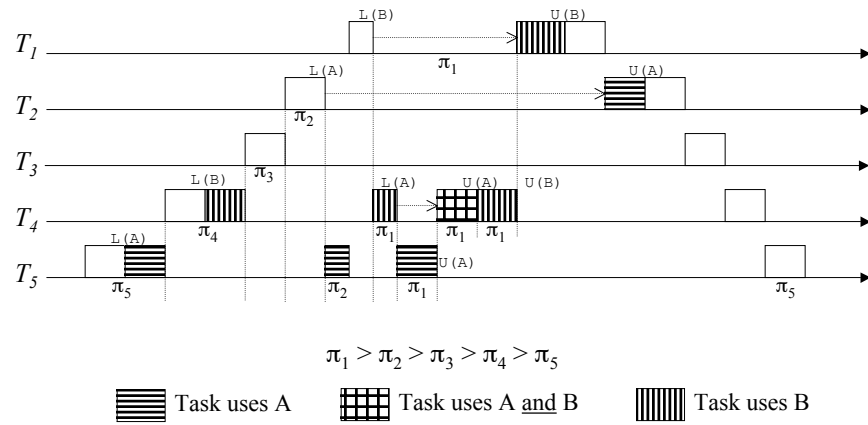
Terminology

- A job is directly blocked when it requests a resource R_i , i.e. executes a $\text{lock}(R_i)$, but no resource of type R_i is available.
- The scheduler grants the lock request, i.e. allocates the requested resource to the job, according to the resource allocation rules, as soon as the resources become available.
- J' directly blocks J if J' holds some resources that J has requested.
- Priority Inheritance:
 - Basic strategy for controlling priority inversion:
Let π be the priority of J
and π' be the priority of J'
and $\pi' < \pi$
then the priority of J' is set to π whenever J' directly blocks J .
- New forms of blocking may be introduced by the resource management policy to control priority inversion and/or prevent deadlocks.

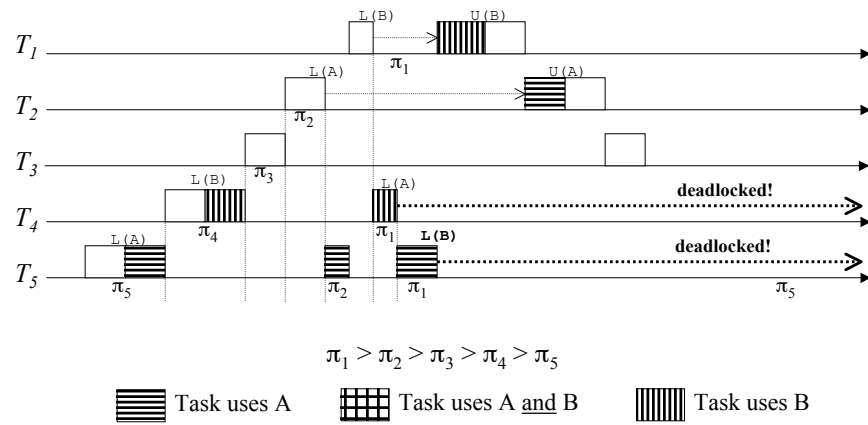
Basic Priority-Inheritance Protocol

- Jobs that are not blocked are scheduled according to a priority-driven algorithm preemptively on a processor.
- Priorities of tasks are fixed, except for the conditions described below:
 - A job J requests a resource R by executing $\text{lock}(R)$
 - If R is available, it is allocated to J . J then continues to execute and releases R by executing $\text{unlock}(R)$
 - If R is allocated to J' , J' directly blocks J . The request for R is denied.
 - However: Let π = priority of J when executing $\text{lock}(R)$
 π' = priority of J' at the same time
 - For as long as J' holds R , its priority is $\max(\pi, \pi')$ and returns to π' when it releases R .
 - That is: J' inherits the priority of J when J' directly blocks J and J has a higher priority.
- Priority Inheritance is transitive.

Example: Priority Inheritance Protocol




Example: Priority Inheritance Protocol



Problem: If T_5 tries to lock (B) while it has priority π_1 , we have a deadlock!

Properties of PIP

- It does not prevent deadlock.
- Task can be blocked directly by a task with a lower priority at most once, for the duration of the (outmost) critical section.
- Consider a task whose priority is higher than n other tasks:



- Each of the lower-priority tasks can directly block the task at most once.
- A task outside the critical section cannot directly block a higher-priority task.